Figure 3 | Patient-level benchmark: is augmentation beneficial?

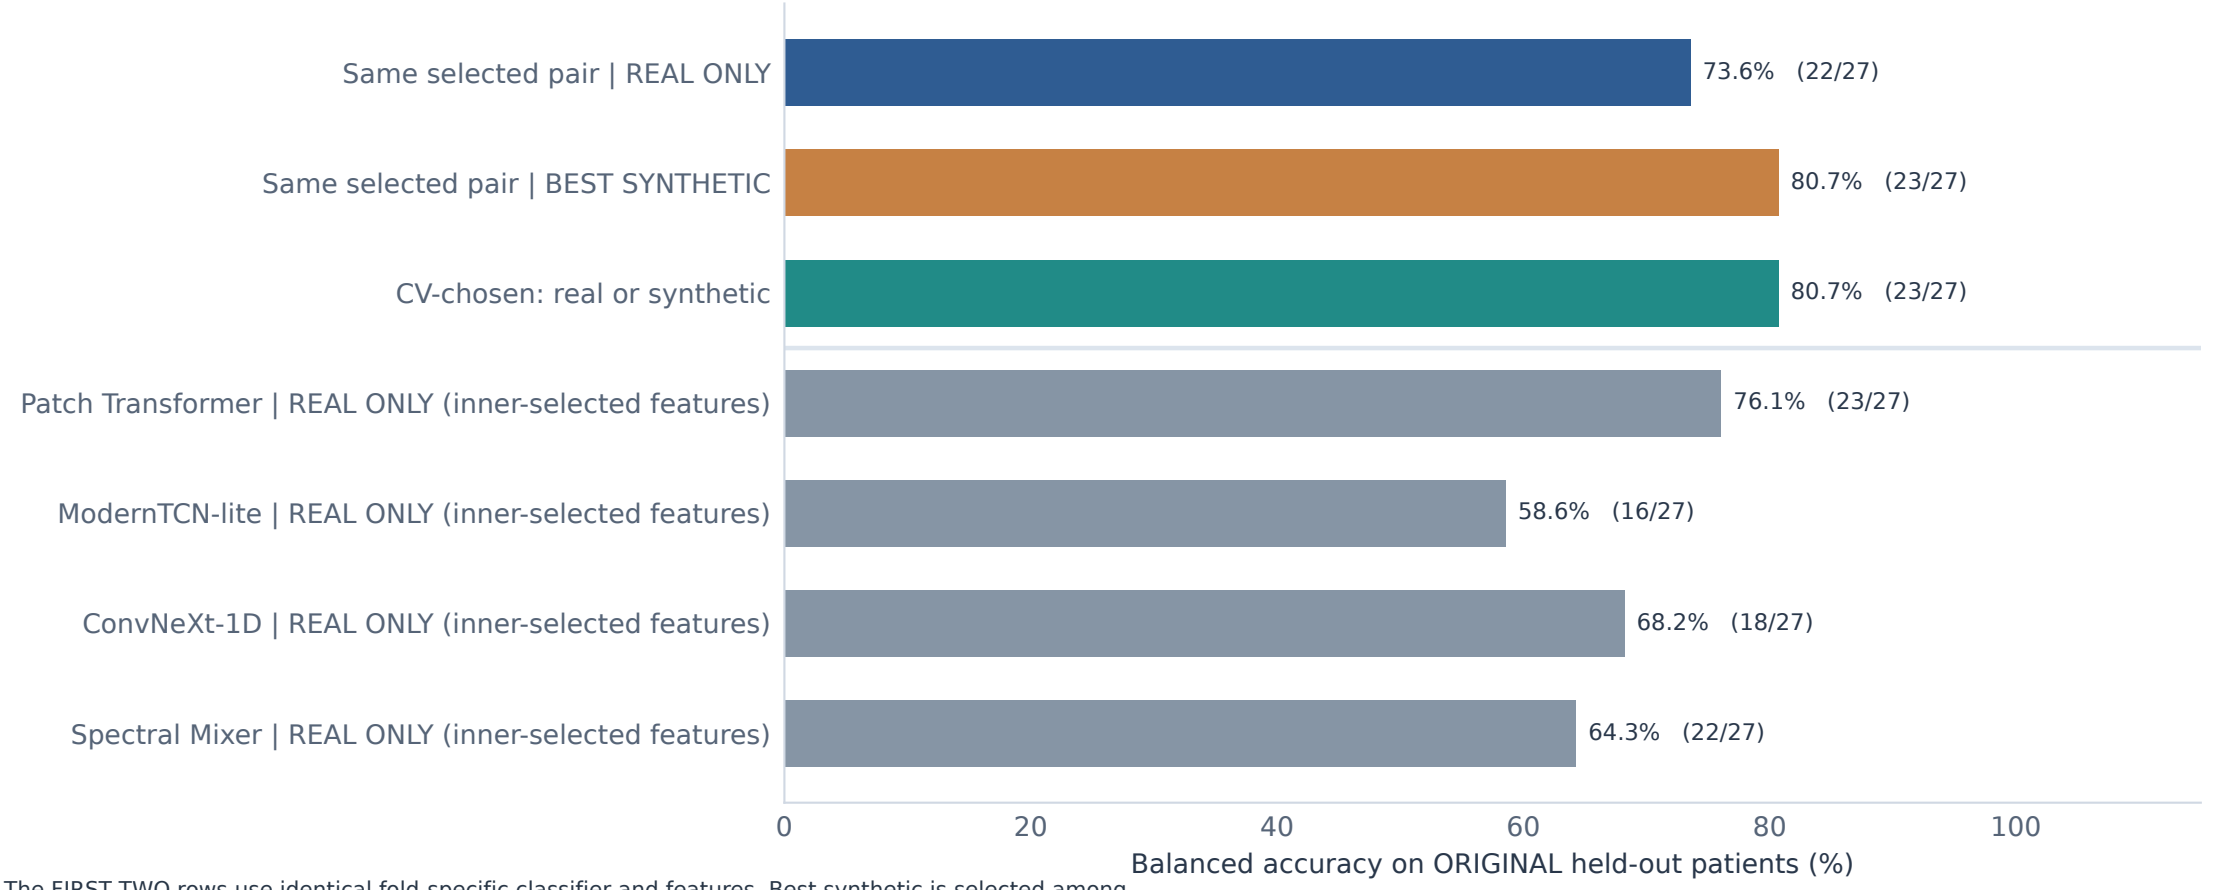

The FIRST TWO rows use identical fold-specific classifier and features. Best synthetic is selected among generators even when real-only wins the probe. Only training data change. All other modern classifiers are REAL-ONLY baselines.